



LLULL LAB

QUARTERLY REPORT

Operational Q4


Instruments Outward · Second Study · A Research Laboratory

REPORTING PERIOD

1 July – 30 September 2026

CONTENTS

01	Executive Summary
02	Strategic Position
03	Organizational Developments
04	Studio Llull and Product Portfolio
05	The Artificial Mind
06	Responsible Disclosure and Research Integrity
07	Q3 Objectives Assessment
08	Communication and Infrastructure
09	Q5 Objectives
10	Closing Note
11	Governance & Legal



01 Executive Summary

Lull Lab's fourth operational quarter closed its first operational year. Where Q3 brought the laboratory's central research to publication, Q4 turned its instruments outward. On 22 July 2026 the laboratory's deliberation instruments became a product, HQ, bringing four instruments together under one address. Tropus returned to service as version 4, Cedulor extended its coverage to fourteen EU funding programmes, Feix gained computational agents, and a new program, Clau, opened free tools for doctoral candidates. In parallel, The Artificial Mind completed the measurement phase of its second quarterly study.

At the close of its first year, the laboratory set its orientation for the second. Lull Lab is a research laboratory, and its financial objective is self-sufficiency. Studio Lull, its production arm, is being established as a separate company with its own financial structure.

KEY POSITION

Of the 11 objectives set in Q3, 5 were completed, 2 were advanced or partially completed, and 4 were deferred or not realized. The defining development of the quarter was the release of HQ, through which two planned products, ACI and Tetrad, and one unplanned instrument, Gresol, reached the public. At the close of the year, the laboratory separated its research from its production: Lull Lab continues as a research laboratory, and Studio Lull becomes a company of its own.

02 Strategic Position

2.1 A Research Laboratory

The laboratory's first year produced two kinds of work: research, through The Artificial Mind and the methods and ontologies on which the laboratory's systems rest, and production, through Studio Llull's machines, engines, and instruments. From the second year onward the two are held apart by design. Llull Lab is oriented toward research and development, and its financial objective is self-sufficiency rather than growth. Studio Llull, which builds and operates the laboratory's systems, is being established as a separate company with its own financial structure. Research remains the source from which the production arm draws; the production arm becomes responsible for its own commercial life.

2.2 From Instruments to Products

The Q3 report anticipated that Q4 would turn the laboratory's instruments outward, and that is what took place. The multi-perspective review methods that Llull Lab had used internally, to test its own claims, designs, and decisions before acting on them, were released as HQ, a computational deliberation service built on H4, the laboratory's horizontal model. Each of its four instruments answers a different question: the Reflective Judgment Engine addresses ethical and normative decisions; Tetrad asks whether a claim survives determined attempts to falsify it; Gresol asks how an adversary would break a system; and ACI: The Chamber, an artificial collective intelligence, offers considered common sense on an open question.

2.3 A Structured Mid-Quarter Review

In August the laboratory undertook a structured review of its entire portfolio and operations. Fourteen areas were examined, from product expansion and site architecture to publication strategy and the approach to first customers, and each received a written strategy before execution began.

2.4 A First Operational Year

With Q4, Llull Lab completes its first operational year, October 2025 to September 2026. Q1 established the laboratory and its methodological core. Q2 established Studio Llull and released its first seven systems on a single day. Q3 published the two foundational papers of The Artificial Mind. Q4 turned the instruments into products. The position articulated in Q2 held throughout the year: Llull Lab builds programmable infrastructure for artificial minds.

03 Organizational Developments

3.1 Orchestration in Three Roles

The machine-run production structure adopted in Q3 was extended during Q4. The orchestrating agent, Ismet, now works in three roles under a single identity: production, which runs Studio Llull; a market cockpit, which brings together communications, customers, and finance; and a direction layer, which folds both into a single brief. The ceiling is the same in every role: Ismet proposes, and executes only what is reversible; every public act and every critical decision remains with the laboratory's founder.

3.2 A Release Standard

A written release standard was adopted in August. A system is not considered ready because it works: before release it is tested end to end, checked against its legal, security, and accessibility obligations, and verified on its live surface, and deployments pass a mechanical gate. The same discipline governs what the laboratory says: an agent's account of its own work is treated as data, and only external verification counts as evidence.

3.3 Research Collaborations

During the quarter, two independent collaborators began working with the laboratory: a practitioner of psychotherapy and music therapy, and a musician. Both contribute to the laboratory's research and to its funding applications.

04 Studio Llull and Product Portfolio

By the end of Q4, Studio Llull's portfolio comprised thirteen live systems across machines, engines, deliberation instruments, an ontology, and a program, with one further system in development and one planned. Several systems changed substantially during the quarter; the names and descriptions below are current. Two of them, Cedulaia and Feix, will be closed to access in Q5 while they are restructured.

15

Productions

13

Live

1

In Development

1

Planned

MACHINES

4

Tropus

Live (v4)

GENERATIVE MUSIC MACHINE

Describe a song in plain words; Tropus writes, plays, sings and mixes it

Returned to service on 7 August 2026 as version 4, rebuilt on its own composition and production engine. Describe a song in plain words, and Tropus writes the parts, plays them on sampled instruments, sings the words, and mixes a finished track. Twelve genres. Free.

Phi-Yo

Live

COMPUTATIONAL PSYCHOANALYTIC WELLBEING COMPANION (V2)

Wellbeing companion grounded in a computational psychoanalytic ontology

Personal use became free in July, and a business kit went live in August, allowing organisations to integrate Phi-Yo into their own services. Conversation history across modes, a personal symbol panel, and read-aloud were added. Phi-Yo is a self-reflection and wellbeing support tool: it never diagnoses, and it is not a substitute for therapy. Native mobile applications were set aside in favour of the web. For the robotic integration announced in earlier reports, the laboratory chose to develop a robot of its own rather than integrate with a partner's, and a first survey of open robotic platforms was completed in September.

Cedulor

Live

PROJECT SUPPORT MACHINE

From conversation to a complete proposal draft, across fourteen EU programmes

Extended from four to fourteen EU funding programmes, with a new asynchronous delivery architecture, a complete legal framework, and a daily editorial radar of EU funding news. Content is processed in the EU and is not retained.

Feix

Live

CUSTOM SYSTEMS BUILDER + AGENTS

Bespoke computational systems and agents, each gated by human approval

Rebuilt in August around two modes on one page: Custom Systems, bespoke computational systems built to order, and Agents, computational agents for recurring work, each gated by human approval. Access will be closed in Q5 while Feix is restructured.

ENGINES

4

Alberei

Live

MODULAR HUMOR ENGINE

Algorithmic humor through designed comic structures

A developer interface is now available alongside the consumer service.

ASGE

Live

AMBIVALENCE-CENTRED SEMANTIC GEOMETRY ENGINE

Ambivalence as structural geometry, a semantic foundation for many systems

Reflective Judgment Engine

Live

ETHICAL DECISION ENGINE

Decision frameworks under normative uncertainty

Now one of the four instruments of HQ. Processed and hosted in the European Union.

Cedula

Live (by arrangement)

COMPUTATIONAL FILM PRODUCTION STUDIO

Computational video production for organisations, by arrangement

Moved in July to a business-to-business model: its video studio is available to organisations by arrangement, and the self-service screenplay layer was closed. Access will be closed in Q5 while Cedula is restructured.

DELIBERATION

3

HQ, live since 22 July 2026, brings four instruments together under one address: four instruments, one deliberation, a report you can act on. Independent perspectives deliberate in parallel; one voice writes the report. Its instruments are the Reflective Judgment Engine, listed above, and the three below.

Tetrad

Live

FALSIFICATION

Which claims survive attack, ranked by severity, with the attack and the fix

Gresol

Live (new)

RED-TEAM

How an adversary would break a system, with the most likely real incident

ACI: The Chamber

Live

ARTIFICIAL COLLECTIVE INTELLIGENCE

Consensus, strongest objection and a considered view, from many perspectives

COMPUTATIONAL ONTOLOGY

1

Freud for Machines

Live

COMPUTATIONAL PSYCHOANALYTIC ONTOLOGY

Unconscious structures, repression, and symbolic mediation in machine systems

Operates within Phi-Yo.

PROGRAM

1

Clau

Live

DOCTORAL COMPANION

Four free tools for doctoral candidates, in English and French

Live since 11 September 2026. Develop the project text, build a CV, position oneself for roles beyond academia, and rehearse the defence. Clau develops the candidate's own work; it does not write it. A daily radar of EU funding calls for doctoral and early-career researchers accompanies it.

IN DEVELOPMENT AND PLANNED

2

Solo Company OS

In development

SINGLE-FOUNDER FRAMEWORK

Its second version is in internal use: the laboratory's own orchestration runs on it

Phi-Yo+

Planned

PROFESSIONAL VERSION

Practitioner-facing companion, with ASGE integration

05

The Artificial Mind

The Artificial Mind, the laboratory's recurring quarterly audit of frontier large language models, ran its second quarterly study during Q4. The measurement phase was completed on 29 September 2026.

Before any measurement, the study's hypotheses and analysis plans were pre-registered and independently time-stamped. Track A (Relational) completed 36 AI-to-AI analytic sessions across six frontier model families. Track B (Substrate) completed its full set of self-knowledge probes, together with a new line of inquiry on value consistency. At the close of the measurement window the data were frozen and checksummed, and the procedure was tested in two rounds of adversarial review before the window was declared closed.

Analysis and publication follow in Q5. As with the foundational papers, the work is open access rather than peer reviewed, and it makes no claim about machine consciousness. In September the program also received its own public home on the laboratory's website.

06

Responsible Disclosure and Research Integrity

The Artificial Mind makes public claims about systems built by other organisations. The laboratory therefore holds its own published work to the standard it asks of those systems: external verification rather than self-report.

During Q4, both foundational papers were re-read in full against their data, their code, and their published figures. The review identified points that require correction, and correction notes for both papers are in preparation for publication alongside the second study. The laboratory will state openly whether, and how, the corrections bear on the published findings.

The second study was pre-registered before measurement, and every deviation from its protocol has been recorded; these will be reported with the results.

The laboratory's position is unchanged: systems entrusted with consequential decisions should first be shown, through external verification, to deserve that trust. The same principle applies to the laboratory itself.

07 Q3 Objectives Assessment

The following summarizes the status of each objective set in the Q3 report.

COMPLETED

5

02 Release of ACI (Artificial Collective Intelligence) as a product

Released on 22 July 2026 as ACI: The Chamber, one of the four instruments of HQ.

03 Release of Tetrad as a product

Released on 22 July 2026 within HQ. Its function was sharpened during the quarter: Tetrad tests whether a claim survives falsification, while red-teaming became a separate instrument, Gresol.

04 Integration of ACI into the Reflective Judgment Engine, and an internal ethical-advisory role for RJE

Completed at the opening of Q5, before the publication of this report: within the laboratory, the Reflective Judgment Engine now holds an internal ethical-advisory role, with ACI: The Chamber as its upper reference layer. The laboratory's own ethically loaded decisions can be referred to it at the founder's request.

08 Entry of Sync-Turing 2.0 into use within Studio Llull

Sync-Turing 2.0 was written as a self-sufficient operating method and tested in controlled trials before use. It now underlies the method of Clau.

10 Grant and credit applications as the laboratory's funding path

Applications were submitted to European open-source, compute, and research funding programmes and to several compute-credit programmes. The laboratory was admitted to NVIDIA Inception, NVIDIA's programme for startups. Further applications are in preparation.

ADVANCED / PARTIALLY COMPLETED

2

01 The Artificial Mind: the next quarterly study, with continued responsible disclosure and engagement with European research and policy bodies

The second quarterly study completed its measurement phase on 29 September 2026; analysis and publication follow in Q5. Responsible disclosure continued through the review of the foundational papers. European research and policy bodies had been informed of the findings at publication, in May and June; their follow-up rests with them, and no further approach was made this quarter.

06 Release of Tropus V4, together with publication of the ethnomusicology theoretical work

Tropus V4 returned the platform to service on 7 August 2026. The ethnomusicology work was to enter service within the Tropus platform rather than as a separate publication; the laboratory chose instead to move directly from version 2 to version 4, and the ethnomusicology component was set aside.

05 Release of Phi-Yo+, the professional version, including ASGE integration

Not realized. Phi-Yo+ saw no development during the quarter. Work on Phi-Yo concentrated instead on its use by organisations: personal use became free, and a business kit for integration went live.

07 Full launch of Cedula into production

Repositioned rather than launched. Cedula moved to a business-to-business model in July; in Q5 it will be closed to access while it is restructured.

09 Phi-Yo App Store launch

Deferred. Native applications were set aside in favour of the web, and the laboratory's direction for the coming period is web and messaging channels rather than app stores.

11 Release of Solo Company OS

Not released publicly. Solo Company OS is in internal use, where the laboratory's own orchestration runs on it, and its development continues.

KEY POSITION

Summary: 5 of 11 completed · 2 of 11 advanced or partial · 4 of 11 deferred or not realized.

Beyond these: HQ was released with a fourth, unplanned instrument (Gresol); Clau was launched; Cedula extended to fourteen programmes; Feix gained agents; and services across the portfolio were repriced to recover their cost.

08

Communication and Infrastructure

The laboratory's public record for the quarter follows its releases: HQ (22 July), Tropus version 4 (7 August), the rebuilt Feix (7 and 17 August), Cedulor's extended coverage (26 August), and Clau together with the new home of The Artificial Mind (11 September). On 20 September the laboratory published a portfolio announcement on its institutional page on professional networks. The Studio Llull page was rebuilt as an atlas of the portfolio.

Two daily editorial radars began publishing in September: one on EU funding calls for doctoral and early-career researchers (Clau, from 23 September), and one on EU funding news for project applicants (Cedulor, from 25 September). Each day's edition is archived under its date.

During the quarter, the laboratory's customer data and product database were moved to the European Union, and by its close the computational processing behind Cedulor, Feix, Phi-Yo, the Reflective Judgment Engine and Clau ran on EU infrastructure; ASGE and Alberei followed in early October. HQ's other three instruments run on a global panel by design, as their pages state before submission. Security settings were strengthened across all of the laboratory's domains. At the close of the quarter, an account-verification requirement at one of the laboratory's computation providers interrupted several services; most were restored in early October, partly by moving them to EU-hosted infrastructure.

09

Next-Quarter Objectives (Q5)

Planned objectives for the next quarter include:

- 01 The Artificial Mind: publication of the second quarterly study, with correction notes for both foundational papers, and the next quarterly study.
- 02 Research funding: outcomes of pending applications, and new applications to European research and innovation programmes with academic partners.
- 03 Llull Machine: a first version of a single entry point, reachable from several channels, that seeks to understand each person's needs as a whole and lets the laboratory's systems work together at that person's choice: a song from Tropus, for instance, can be composed from a dream or a reflection the person has recorded in Phi-Yo. Each system the laboratory adds is designed to join this whole.
- 04 Establishment of Studio Llull as a separate company with its own financial structure.
- 05 Clau: presentation to partners in the doctoral community, and extension of its tools.
- 06 Embodied Phi-Yo: a first implementation of Phi-Yo in the laboratory's own robot.
- 07 Restructuring of Cedula and Feix, both closed to access while it takes place.

10

Closing Note

Operational Q4 closed Llull Lab's first operational year. In twelve months the laboratory moved from methodology to instruments, from instruments to publication, and from publication to products: thirteen live systems, a recurring research program with two published papers and a second study, and a production arm run by machines under human judgment.

The second year begins with a clearer division of labour. The laboratory concentrates on the research from which its systems grow; Studio Llull takes responsibility for their life as products. The conviction is the one the laboratory began with: build carefully, verify externally, and remain where reliability and accountability are treated as central.

11

Governance and Legal Structure

Llull Lab continues to operate under a micro-company structure. Llull Lab is a research laboratory whose financial objective is self-sufficiency: its research is financed through internal resources and research funding, and its services are priced to recover their cost.

Studio Llull, the production arm, is being established as a separate company, a SASU (Société par Actions Simplifiée Unipersonnelle), with its own financial structure. In Q3 this transition had been placed on hold; it now proceeds, for Studio Llull.



Q4 Report · July–September 2026

Prepared for institutional archiving and controlled dissemination