



LLULL LAB

QUARTERLY REPORT

Operational Q3


The Artificial Mind · Open-Access Publication · Machine-Run Studio

REPORTING PERIOD

1 April – 30 June 2026

CONTENTS

01	Executive Summary
02	Strategic Position
03	Organizational Developments
04	Studio Llull and Product Portfolio
05	The Artificial Mind
06	Responsible Disclosure and Strategic Positioning
07	Q2 Objectives Assessment
08	Communication and Infrastructure
09	Financial Position
10	Q4 Objectives
11	Closing Note
12	Governance & Legal



01 Executive Summary

Lull Lab's third operational quarter was the quarter in which the laboratory's core research reached publication. On 21 June 2026, the two foundational papers of The Artificial Mind, the laboratory's recurring quarterly audit of frontier large language models (LLMs), were released open access. The program proceeds through two complementary methods: AI-to-AI psychoanalysis (Track A) and substrate self-knowledge probes (Track B). In parallel, Studio Lull completed its transition to a fully machine-run production structure.

KEY POSITION

Of the 13 objectives set in Q2, 5 were completed, 4 were advanced or partially completed, and 4 were deliberately deferred. The defining achievement of the quarter was the open-access release of the two foundational papers of The Artificial Mind. Studio Lull completed its transition to a fully machine-run production structure, and the laboratory's research focus took precedence over commercial development. Beyond these objectives, Cedula gained computational vision-language understanding, and The Artificial Mind was unified into a single recurring program.

All Q3 activities were fully financed through internal resources. No external funding or operating revenue was recorded during this period. As in Q1 and Q2, this remains a deliberate choice rather than a constraint.

02 Strategic Position

2.1 The Artificial Mind as a Research Program

The research nucleus established in Q1 as a dedicated unit has matured into a recurring research program: The Artificial Mind, the laboratory's quarterly audit of frontier large language models. The program proceeds through two complementary methods: AI-to-AI psychoanalysis, which examines how an artificial mind conducts and undergoes the analytic relationship, and substrate self-knowledge probes, which test what such a system can reliably report about itself. Rather than a fixed organizational compartment, it now operates as a continuing line of inquiry with a quarterly cadence. All of the laboratory's instruments, engines, and ontologies remain conceptually anchored within this program.

2.2 Vision Continuity

The laboratory's position, consolidated in Q2, held throughout Q3: Llull Lab builds programmable infrastructure for artificial minds. Studio Llull, its production arm, builds the systems in which capacities such as affect, ethical judgment, humor, dreaming, and consciousness can be composed, orchestrated, and studied. Q3 deepened this position rather than broadening it, with the laboratory's energy concentrated on bringing its central research to publication.

03

Organizational Developments

3.1 Transition to a Machine-Run Production Structure

During Q3, Llull Lab once more tested human collaboration as a basis for Studio Llull's production work. This approach again did not prove the right fit for the laboratory's needs, and the decision was taken to hand the production of Studio Llull entirely to machine management. The orchestrating agent, Ismet, was brought into the core team.

This step did not introduce a new principle so much as it clarified one. The zero-employee, machine-native model adopted in Q2 is now fully realized in practice: Studio Llull is designed, built, and operated by AI agents. For tasks that machines cannot yet perform, human labour is engaged temporarily. What began as a response to the difficulties of human resource management has become a settled architectural principle. The laboratory builds systems for artificial minds, and it operates through them.

3.2 Neurodivergent Collaboration

The open call for project-based collaboration with neurodivergent individuals, announced in Q2, produced its first results during Q3: neurodivergent collaborators began contributing to the laboratory's project-based work, and further individuals have since expressed interest in taking part. This reaffirms a core institutional conviction, that neurodivergent minds are not merely welcome but valuable at the intersection of computation, cognition, and culture.

04 Studio Llull and Product Portfolio

Studio Llull is Llull Lab's production arm. By the end of Q3 its established portfolio comprised nine productions: eight live, and one (Tropus) in maintenance. Several further systems were in development for release in Q4. Several productions were repositioned during the quarter; the names and descriptions below are current.

9

Productions

8

Live

1

In Maintenance

MACHINES

4

Tropus

Offline (maintenance)

GENERATIVE MUSIC MACHINE

Musical generation and conceptual inquiry in a dual-mode architecture

Tropus remained offline through Q3, having entered a major infrastructure transition on 1 March 2026. During this period the platform's engine was substantially rebuilt, including new vocal synthesis and per-stem mix handling. Version 4 is in development, with release planned for Q4.

Phi-Yo

Live

COMPUTATIONAL PSYCHOANALYTIC WELLBEING COMPANION (V2)

Wellbeing companion grounded in a computational psychoanalytic ontology

Repositioned during Q3 from a dream analysis tool into a wellbeing companion. The Freud for Machines ontology continues to operate within Phi-Yo as its interpretative foundation. Native mobile applications are nearing release: the iOS version is ready, and the Android version is in testing.

Cedulor

Live

EU GRANT PROPOSAL MACHINE

Conversation to a complete EU proposal document, across four programmes

Expanded during Q3. Where its scope was previously limited to Horizon Europe, Cedulor now supports four EU funding programmes (Horizon Europe, MSCA, ERC, and EIC Accelerator), each with a dedicated pipeline.

Feix

Live

CUSTOM SYSTEMS BUILDER

Bespoke operational systems, delivered from brief to live product

Positioning refined during the quarter to bespoke operational systems, delivered from brief to live product.

ENGINES

4

Alberei

Live

MODULAR HUMOR ENGINE

Algorithmic humor through designed comic structures

ASGE

Live

AMBIVALENCE-CENTRED SEMANTIC GEOMETRY ENGINE

Ambivalence as structural geometry, a semantic foundation for many systems

Reflective Judgment Engine

Live

ETHICAL DECISION ENGINE

Decision frameworks under normative uncertainty

Cedula

Live (screenplay)

COMPUTATIONAL FILM PRODUCTION STUDIO

Story to screenplay, the first layer of a full film-production pipeline

The story-to-screenplay layer is live. During Q3, computational vision-language understanding was added to the system. The full film-production pipeline is in development, with a full launch planned for Q4.

COMPUTATIONAL ONTOLOGY

1

Freud for Machines

Live

COMPUTATIONAL PSYCHOANALYTIC ONTOLOGY

Unconscious structures, repression, and symbolic mediation in machine systems

Operates as an integrated component within Phi-Yo.

IN DEVELOPMENT

4

ACI

In development

ARTIFICIAL COLLECTIVE INTELLIGENCE

Collective-reasoning instrument, a relational alternative to AGI

Tetrad

In development

MULTI-PERSPECTIVE REVIEW BOARD

Multi-perspective review of code and design

Phi-Yo+

In development

PROFESSIONAL VERSION

Practitioner-facing companion, with ASGE integration

Solo Company OS

In development

SINGLE-FOUNDER BOOTSTRAP FRAMEWORK

Operating framework for solo companies

05

The Artificial Mind

Q3 saw the publication of the two foundational papers of The Artificial Mind, the laboratory's recurring quarterly audit of frontier large language models. Both were released open access on 21 June 2026. The program proceeds through two complementary methods, one per paper.

Track A (Relational): "The Analyst Has No Body Either."

DOI 10.5281/zenodo.20779122. A full psychoanalytic-dyad study across six frontier model families, examining subjectivity as an effect of language rather than as a claim about consciousness.

Track B (Substrate): "Basins, Bargains, and the Limits of Self-Report."

DOI 10.5281/zenodo.20779376. A behavioural audit of the same model families, distinguishing the self-claims that can be verified by external probing from those that structurally cannot.

The work is open access rather than peer reviewed, and it makes no claim about machine consciousness. Its central public finding is a structural tendency, observed across multiple frontier model families, for these systems to judge their own reliability more favourably than independent testing supports. The magnitude of this effect is not yet established.

The two papers constitute a frozen baseline. The Artificial Mind is structured as a recurring program: subsequent quarters extend the instrument and revisit the findings on a regular cadence.

06 Responsible Disclosure and Strategic Positioning

The central finding described above, the cross-model tendency toward favourable self-assessment, carries weight well beyond any single product. Its significance is public, not commercial.

Large language models are increasingly placed in high-consequence roles, including reported integration into defence and weapons systems, often before independent reliability testing is complete and on the assumption that such systems can be trusted to assess themselves. Our findings indicate that this assumption cannot be taken for granted, and that a gap of this kind, in settings where decisions may be irreversible, can carry risk not only to human life but to the living world and the planet we share, on a scale that warrants caution.

The laboratory states this plainly and without accusation: systems entrusted with irreversible decisions should first be shown, through external verification, to deserve that trust. In this spirit, Llull Lab informed a range of research, editorial, and institutional bodies of its findings during the quarter.

The quarter also affirmed the laboratory's conviction in remaining EU-based and in contributing to credible European alternatives in foundational AI, where reliability, accountability, and human safety are treated as central rather than incidental.

07

Q2 Objectives Assessment

The following summarizes the status of each objective set in the Q2 report.

COMPLETED

5

01 Entry into machine-to-machine psychoanalysis

Completed and exceeded. The two foundational papers of The Artificial Mind were released open access. The work was established as a recurring research program rather than a single output.

03 Theoretical work for Tropus V4 (ethnomusicology)

The theoretical work is prepared and ready for publication, which will follow in Q4.

06 Expansion of Cedulor

Cedulor was expanded from Horizon Europe alone to four EU funding programmes.

12 Consolidation of the zero-employee operational model

Completed and clarified. Studio Lull's production was handed entirely to machine management, and the orchestrating agent was brought into the core team.

13 General visibility and outreach expansion

Completed. A multilingual institutional outreach effort accompanied the publication, an institutional presence on professional networks was established, and the public-facing surfaces were updated.

ADVANCED / PARTIALLY COMPLETED

3

05 Transformation of Phi-Yo into a wellbeing companion, with ASGE integration

The repositioning of Phi-Yo into a wellbeing companion was completed. ASGE integration has been moved into the forthcoming professional version, Phi-Yo+.

07 Full launch of Cedulia

The screenplay layer is live, and computational vision-language understanding was added during the quarter. The full launch is planned for Q4.

08 Initial deployment of RJE, Alberei, and Phi-Yo in physical robotic systems

The integration layer connecting the Reflective Judgment Engine, Alberei, and Phi-Yo is complete and deployment-ready. First physical deployment is proceeding through partner integration.

IN PREPARATION

1

11 Design of SaaS and API-based revenue models

Preparatory work was begun. It was not prioritized this quarter as research took precedence.

02 Release of Tropus V3

Superseded by Version 4. The platform's engine was substantially rebuilt during the infrastructure transition; release is planned for Q4.

04 Release of H4

Deferred owing to the concentration of effort on the quarter's primary research.

09 Completion of SASU transition

Placed on hold. Given the internal-resource model and the priority on research, the transition was not a current priority.

10 Initiation of the seed funding process

Deferred. Commercial development was not prioritized this quarter.

KEY POSITION

Summary: 5 of 13 completed · 4 of 13 advanced or partial · 4 of 13 deferred.

Beyond these: Cedula gained computational vision-language understanding, and The Artificial Mind was unified into a single recurring program.

08 Communication and Infrastructure

The laboratory's public-facing surfaces were updated during Q3 to reflect the publication of The Artificial Mind and the current state of the production portfolio.

An institutional presence was established on professional networks, extending the laboratory's outreach beyond its existing channels. A dedicated visual series accompanied the publication.

During the quarter, the laboratory held discussions concerning the integration of Phi-Yo into physical robotic systems. These discussions are ongoing.

The Phi-Yo native mobile applications advanced toward release: the iOS version is ready, and the Android version is in testing.

09 Financial Position

All Q3 activities were fully financed through internal resources. No external funding or operating revenue was recorded during this period.

This remains a deliberate choice. With the laboratory's effort concentrated on bringing The Artificial Mind to publication, commercial development was not prioritized during the quarter, and the same orientation is expected to continue into Q4. The laboratory elected to sustain its current rhythm from internal resources, focused on the rapid advance of large language models and robotics, without acceleration from outside capital.

10

Next-Quarter Objectives (Q4)

Planned objectives for the next quarter include:

- 01 The Artificial Mind: the next quarterly study, together with continued responsible disclosure and engagement with European research and policy bodies.
- 02 Release of ACI (Artificial Collective Intelligence) as a product: a collective-reasoning instrument that draws judgment from many independent perspectives. The laboratory frames it as a relational, external alternative to the pursuit of artificial general intelligence (AGI): intelligence composed across many minds rather than concentrated within one.
- 03 Release of Tetrad as a product: a multi-perspective review board for code and design.
- 04 Integration of ACI into the Reflective Judgment Engine as an upper reference layer, and the assignment of an internal ethical-advisory role to the Reflective Judgment Engine.
- 05 Release of Phi-Yo+, the professional version, including ASGE integration.
- 06 Release of Tropus V4, returning the platform to service, together with publication of the ethnomusicology theoretical work.
- 07 Full launch of Cedula into production.
- 08 Entry of Sync-Turing 2.0 into use within Studio Llull.
- 09 Phi-Yo App Store launch (iOS submission and Android release).
- 10 Grant and credit applications as the laboratory's funding path.
- 11 Release of Solo Company OS.

11

Closing Note

Operational Q3 was the quarter in which Llull Lab's central research reached the public. The two foundational papers of The Artificial Mind were released open access, and the laboratory acted on its findings by informing research, editorial, and institutional bodies. Studio Llull completed its transition to a fully machine-run production structure, and the laboratory sustained its work entirely from internal resources.

Q4 will turn the laboratory's instruments outward. Deliberation systems become products, the professional and embodied extensions of existing tools take shape, and the music platform returns to service. The laboratory continues to hold that the most responsible position, in a period of rapid and high-consequence change, is to build carefully, to verify externally, and to remain where reliability and accountability are treated as central.

12

Governance and Legal Structure

Llull Lab continues to operate under a micro-company structure. Preparations for a transition to a SASU (Société par Actions Simplifiée Unipersonnelle) are complete, but the transition, together with any search for outside investment, has been placed on hold. Given the internal-resource model and the priority on research, neither is a current priority; the laboratory will revisit both when its circumstances warrant. Studio Llull will not acquire autonomous company status at this stage; it continues under Llull Lab.



Q3 Report · April–June 2026

Prepared for institutional archiving and controlled dissemination