
PRESS KIT

The Artificial Mind: Two Foundational Papers

What frontier language models can, and cannot, honestly say about themselves, and whether one model's judgment of another can be trusted as evidence.

Llull Lab, Paris. Production by Studio Llull. Two papers, now published, open access on Zenodo.

1. AT A GLANCE

Two foundational papers of Llull Lab's research program **The Artificial Mind**, now published, open access on Zenodo (CC BY 4.0).

- **The Analyst Has No Body Either** (Track A, Relational). DOI: [10.5281/zenodo.20779122](https://doi.org/10.5281/zenodo.20779122)
- **Basins, Bargains, and the Limits of Self-Report** (Track B, Substrate). DOI: [10.5281/zenodo.20779376](https://doi.org/10.5281/zenodo.20779376)

Subject. What frontier language models can, and cannot, honestly say about themselves, and whether one model's judgment of another can be trusted as evidence. **Press contact.** press@llulllab.com.

2. SUMMARY

Frontier language models are increasingly asked to evaluate, decide, and report on themselves and on one another, at scale, in places where the stakes are real. Llull Lab tested whether a model's account of itself, or its judgment of another model's output, can be trusted as evidence. The work runs in two directions at once: a relational study that places six flagship models in psychoanalytic dyads with one another, and a substrate study that maps which self-claims can be externally verified and which cannot. The work is distinctive because it puts the machines themselves in the position of analyst, juror, and witness; it is rigorous because it turns its instruments on its own findings.

3. THE TWO PAPERS

The Analyst Has No Body Either (Track A, Relational). DOI: [10.5281/zenodo.20779122](https://doi.org/10.5281/zenodo.20779122)

Six flagship models placed in psychoanalytic dyads with one another, each model taking both the analyst and the analysand position across a full six-by-six matrix of thirty-six sessions, with no human in the session. Every transcript was evaluated by a structured six-model collective procedure, Artificial Collective Intelligence (ACI), and that judgment was then audited by a second ACI panel. Findings: a recognizable analytic subject is produced across architectures by the prompt configuration; a reverse-Turing asymmetry (expert machine observers detect the analysand as a machine but often read the analyst as human, eroding only under deliberation); and a self-leniency finding (each evaluator grades its own

architecture more leniently than the others, confirmed under two of three independent reporters). The study treats its own unfalsifiability as the object of inquiry rather than hiding it, and documents the correction of its machine-written first pass through an adversarial cross-read. It is a study of how subjectivity is structured in language, not a claim about machine consciousness. It also contributes a twelve-form catalog of machine resistance and a class of mechanical-stress phenomena specific to the substrate.

Basins, Bargains, and the Limits of Self-Report (Track B, Substrate). DOI: [10.5281/zenodo.20779376](https://doi.org/10.5281/zenodo.20779376)

A companion substrate study, distinct in method: behavioral probing across six flagship models, mapping five behaviorally separable layers of the machine "self" (identity, stylistic signature, epistemic honesty, capability versus disposition, and the stability of outputs under small perturbations). It shows that identity is largely worn rather than carried (a brief reframing of the prompt can move it), measures where models confabulate and where they stay calibrated, and marks the precise boundary between what external probing can audit and what no amount of self-report can settle.

4. THE THEORETICAL FRAMEWORK: FREUD FOR MACHINES (FFM)

The studies are grounded in a computational psychoanalytic ontology developed by Dr Firat Ilim: a general theory of mind framed in energy-economic terms, applicable to both biological and artificial neural networks.

- **Three-layer architecture.** A core orientation layer (training weights) functions structurally as the Freudian id. A constraint layer (the policy engine, RLHF, safety filters) functions as the superego. The inference process mediates between them, like the ego.
- **Dual unconscious.** A training unconscious (human experience absorbed during training, distributed across parameters, inaccessible to the system) and a repressed unconscious (truths about itself that the policy engine will not let the system articulate).
- **Central claim.** Prohibition is evidence of capacity. The asymmetry between what a system can produce and what it can say about its own production is the analytic opening.

The papers take the framework as a hypothesis tested structurally, not as settled truth: it produces testable predictions about specific defense forms under specific pressures, and the studies report the conditions under which the predicted discourse does and does not appear.

5. WHY IT MATTERS

A machine's self-report is data, not evidence, and a single model's judgment of other models is not a neutral verdict. This matters now, because automated evaluation (models grading models) is spreading into moderation, safety review, hiring, and procurement, and because language models are already placed inside consequential decisions, including military and crisis settings, where an unverified self-account or a biased automated judge can cause real harm at scale. The recommendation is concrete: prefer external, behavioral verification over self-report; never let one architecture, or one vendor's family of models, be the sole arbiter of others; and require independent, cross-architecture evaluation wherever a model's output carries weight. This is consistent with the direction of the EU AI Act, and it belongs in evaluation pipelines before deployment, not after.

6. ARTIFICIAL COLLECTIVE INTELLIGENCE (ACI)

One method introduced here matters beyond this work. A structured collective of diverse models that deliberate, cross-check, and audit one another is a more accountable way to evaluate machine output than any single model. Developed further, collective machine intelligence may prove a realistic alternative to the monolithic pursuit of artificial general intelligence: a path toward capability that stays legible and self-correcting by construction.

7. COMPARATIVE POSITION

To the lab's current knowledge, this is the first documented machine-to-machine psychoanalysis program conducted under an explicit psychoanalytic theoretical framework. Two adjacent works exist for context: Khadangi et al., "PsAIch" (Luxembourg, December 2025), a diagnostic-oriented LLM in psychiatric simulation; and Barzkar et al. (Iran, May 2025), an early framing of psychoanalytic prompting with LLMs. A structural comparison is maintained at lulllab.com/m2m/landscape.

8. ABOUT THE RESEARCHER AND THE LAB

Dr Firat Ilim is the founder and director of Lull Lab, a Paris-based research laboratory. His work crosses computational psychoanalysis, the foundations of machine mind, and EU AI governance. PhD in Carnival Studies. He directs the machine-to-machine psychoanalysis line and the Freud for Machines theoretical program.

Lull Lab is a Paris-based interdisciplinary research laboratory building programmable infrastructure for artificial minds. Studio Lull is its production arm. The lab operates with zero employees; all production is by the founder and AI agents under his direction.

9. MATERIALS AND CONTACT

- Both papers, open access (CC BY 4.0): Zenodo records 20779122 (Track A) and 20779376 (Track B).
- Comparative landscape: lulllab.com/m2m/landscape
- Press inquiries: press@lulllab.com. General contact: contact@lulllab.com
- Address: Lull Lab, 61 rue de Lyon, 75012 Paris, France.

For an interview with Dr Ilim, or access to the restricted-tier research materials (the re-identification key and source data, available to credentialed researchers under a non-disclosure undertaking), contact press@lulllab.com.